

Deep Reinforcement Learning

Deep Reinforcement Learning


Using deep learning to improve reinforcement learning

Deep learning: Training (deep) neural networks to accurately map inputs to outputs

Reinforcement learning: Learning from experience


Deep Reinforcement Learning

Typical RL scenario



Deep Reinforcement Learning

Typical RL scenario



$$\pi(s) : \mathcal{S} \rightarrow \mathcal{A}$$

Q-Learning

$$\begin{aligned} Q^\pi(s, a) &= E_\pi \{R_t | s_t = s, a_t = a\} \\ &= E_\pi \left\{ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t = s, a_t = a \right\} \\ &= r(s, a) + \gamma Q^\pi(s', a') \end{aligned}$$

Q-Learning

$$\begin{aligned} Q^\pi(s, a) &= E_\pi \{ R_t | s_t = s, a_t = a \} \\ &= E_\pi \left\{ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t = s, a_t = a \right\} \\ &= r(s, a) + \gamma Q^\pi(s', a') \end{aligned}$$

$$\pi^*(s) = \arg \max_a Q(s, a)$$

Q-Learning

$$\begin{aligned} Q^\pi(s, a) &= E_\pi \{ R_t | s_t = s, a_t = a \} \\ &= E_\pi \left\{ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s, a_t = a \right\} \\ &= r(s, a) + \gamma Q^\pi(s', a') \end{aligned}$$

$$\pi^*(s) = \arg \max_a Q(s, a)$$

$$Q(s, a) = r(s, a) + \gamma \max_{a'} Q(s', a')$$

Q-Learning

$$Q(s, a) = r(s, a) + \gamma \max_{a'} Q(s', a')$$

If \mathcal{S} and \mathcal{A} are finite we can simply store the q-value for each (s, a) pair in a lookup table.


Q-Learning

$$Q(s, a) = r(s, a) + \gamma \max_{a'} Q(s', a')$$

If \mathcal{S} and \mathcal{A} are finite we can simply store the q-value for each (s, a) pair in a lookup table.

But if either of them are infinite, or if there are simply too many states (or actions) to store the q-values in a lookup table, we need an alternative.

Deep Q-Learning



Q-Learning

Universal Approximation Theorem

A feed-forward neural network with a single hidden layer containing a finite number of neurons can approximate continuous functions on compact subsets of R^n , under mild assumptions on the activation function.

In other words:

Simple neural networks can represent a wide variety of interesting functions...

Q-Learning

Universal Approximation Theorem

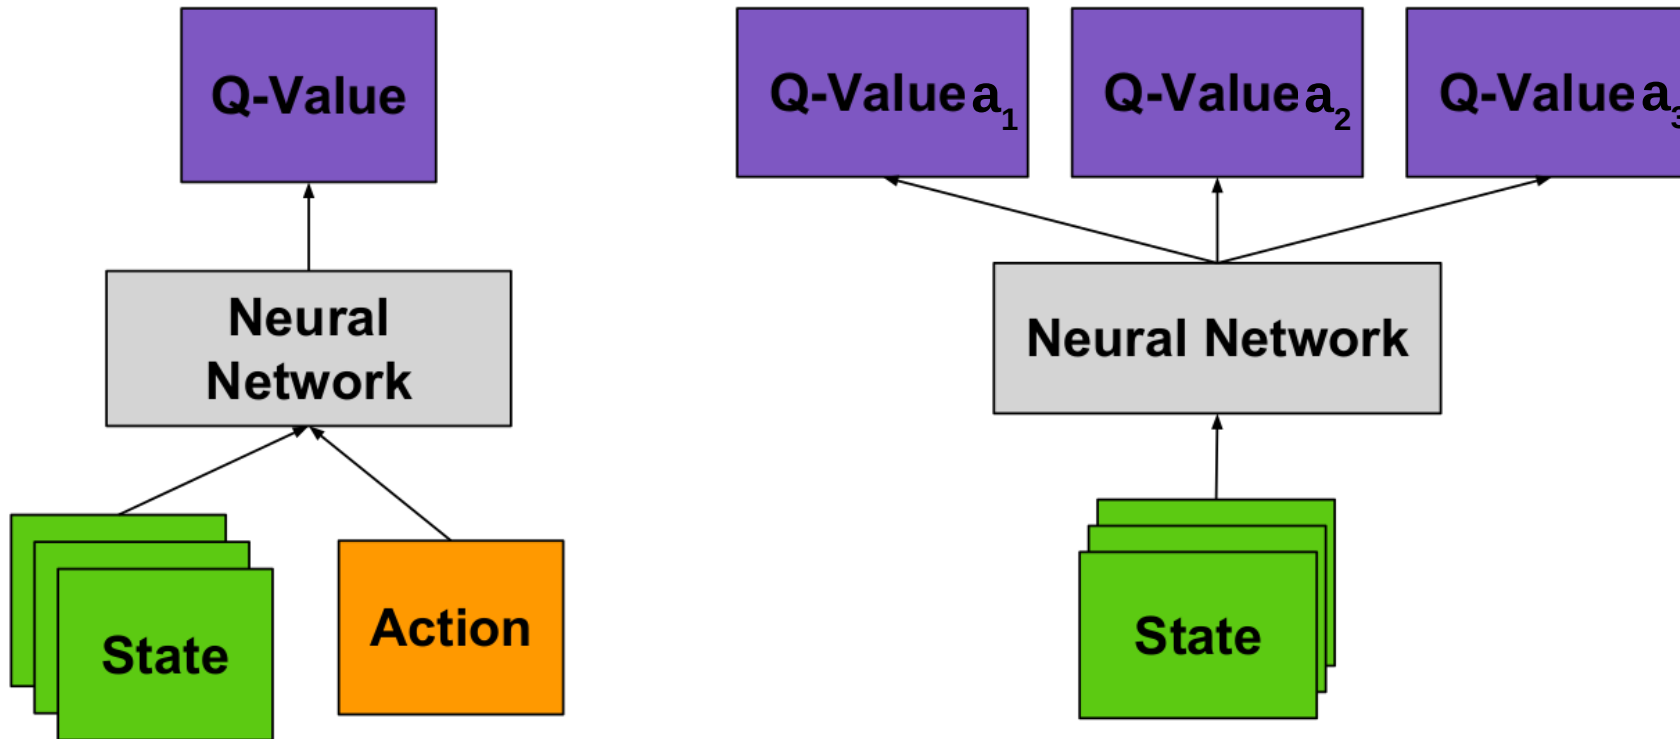
A feed-forward neural network with a single hidden layer containing a finite number of neurons can approximate continuous functions on compact subsets of R^n , under mild assumptions on the activation function.

In other words:

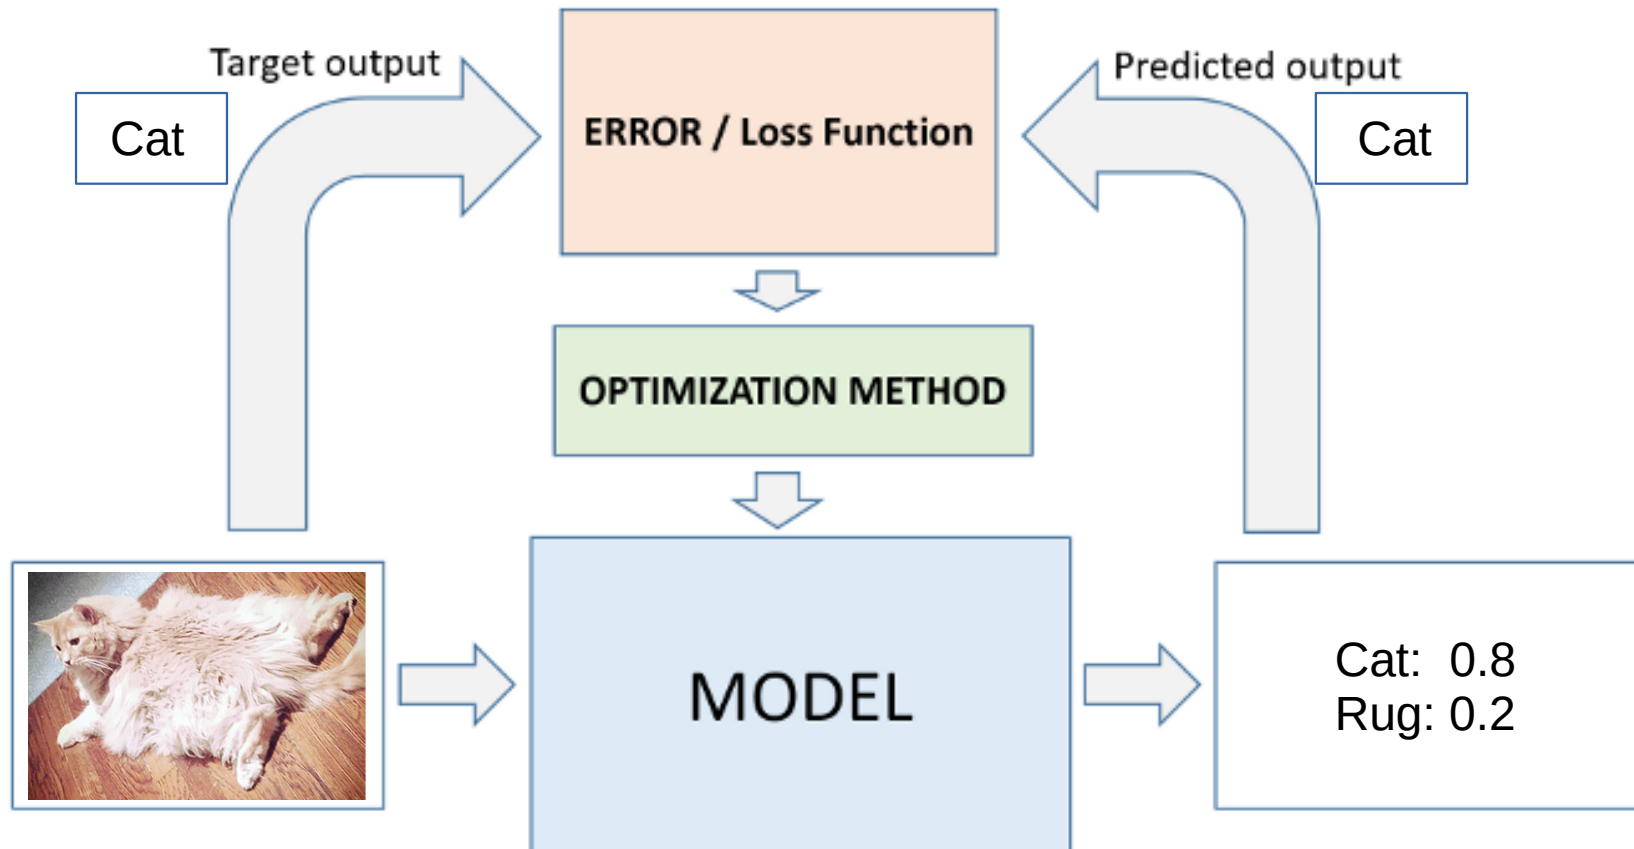
Simple neural networks can represent a wide variety of interesting functions...

... Including Q functions!

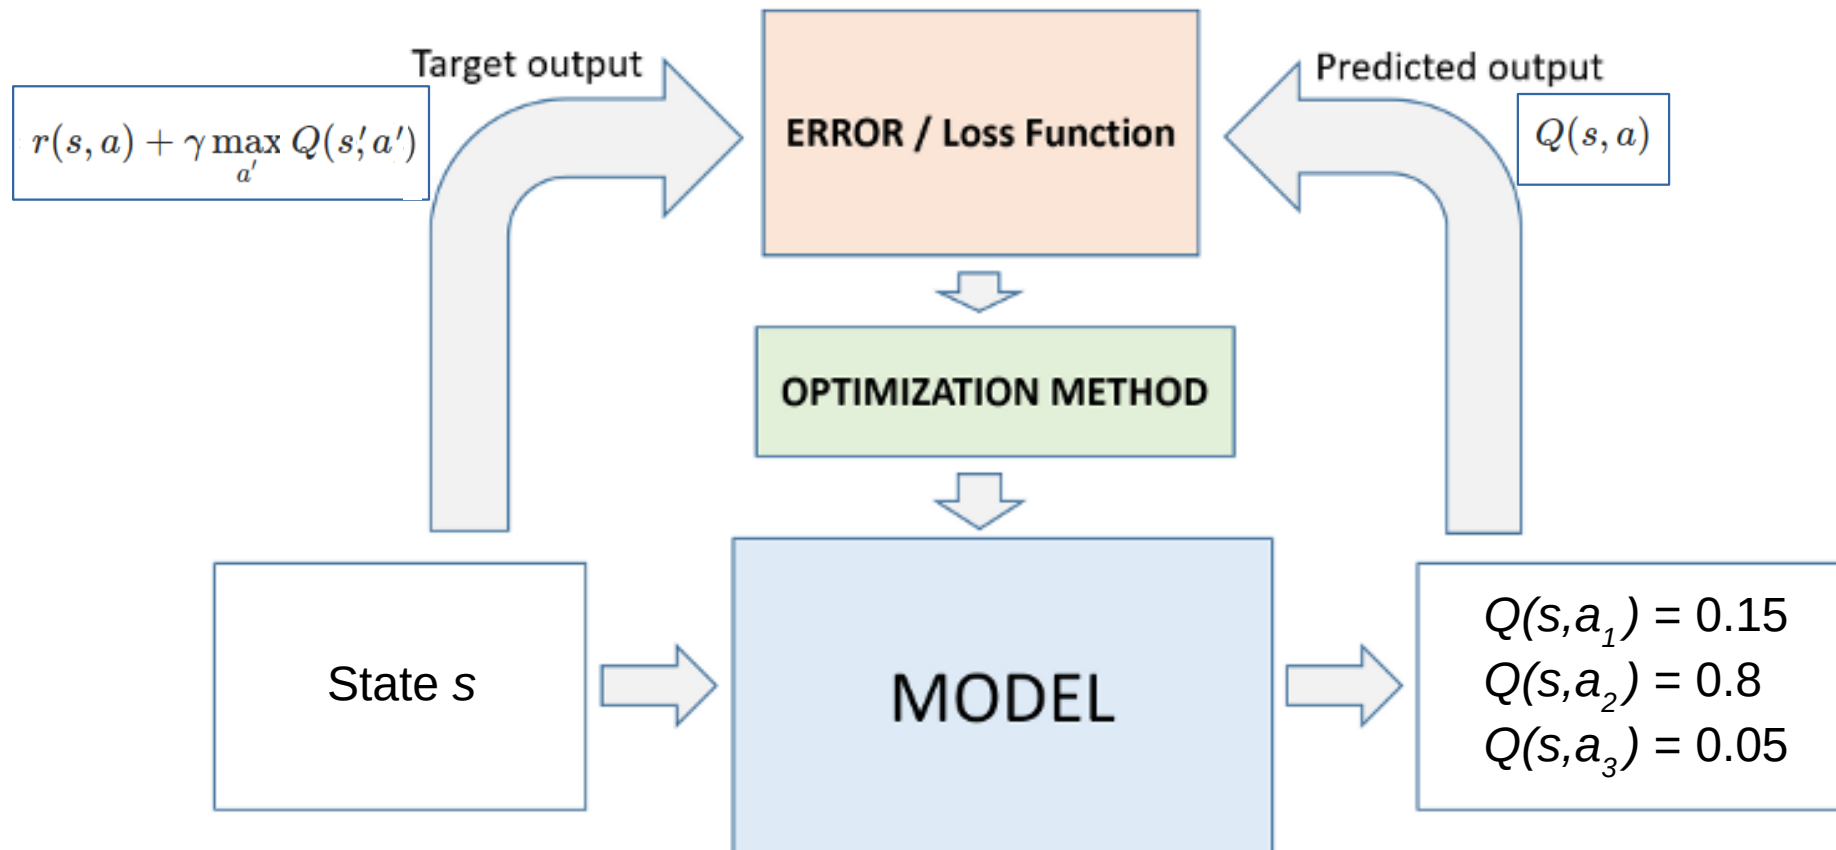
Deep Q-Learning



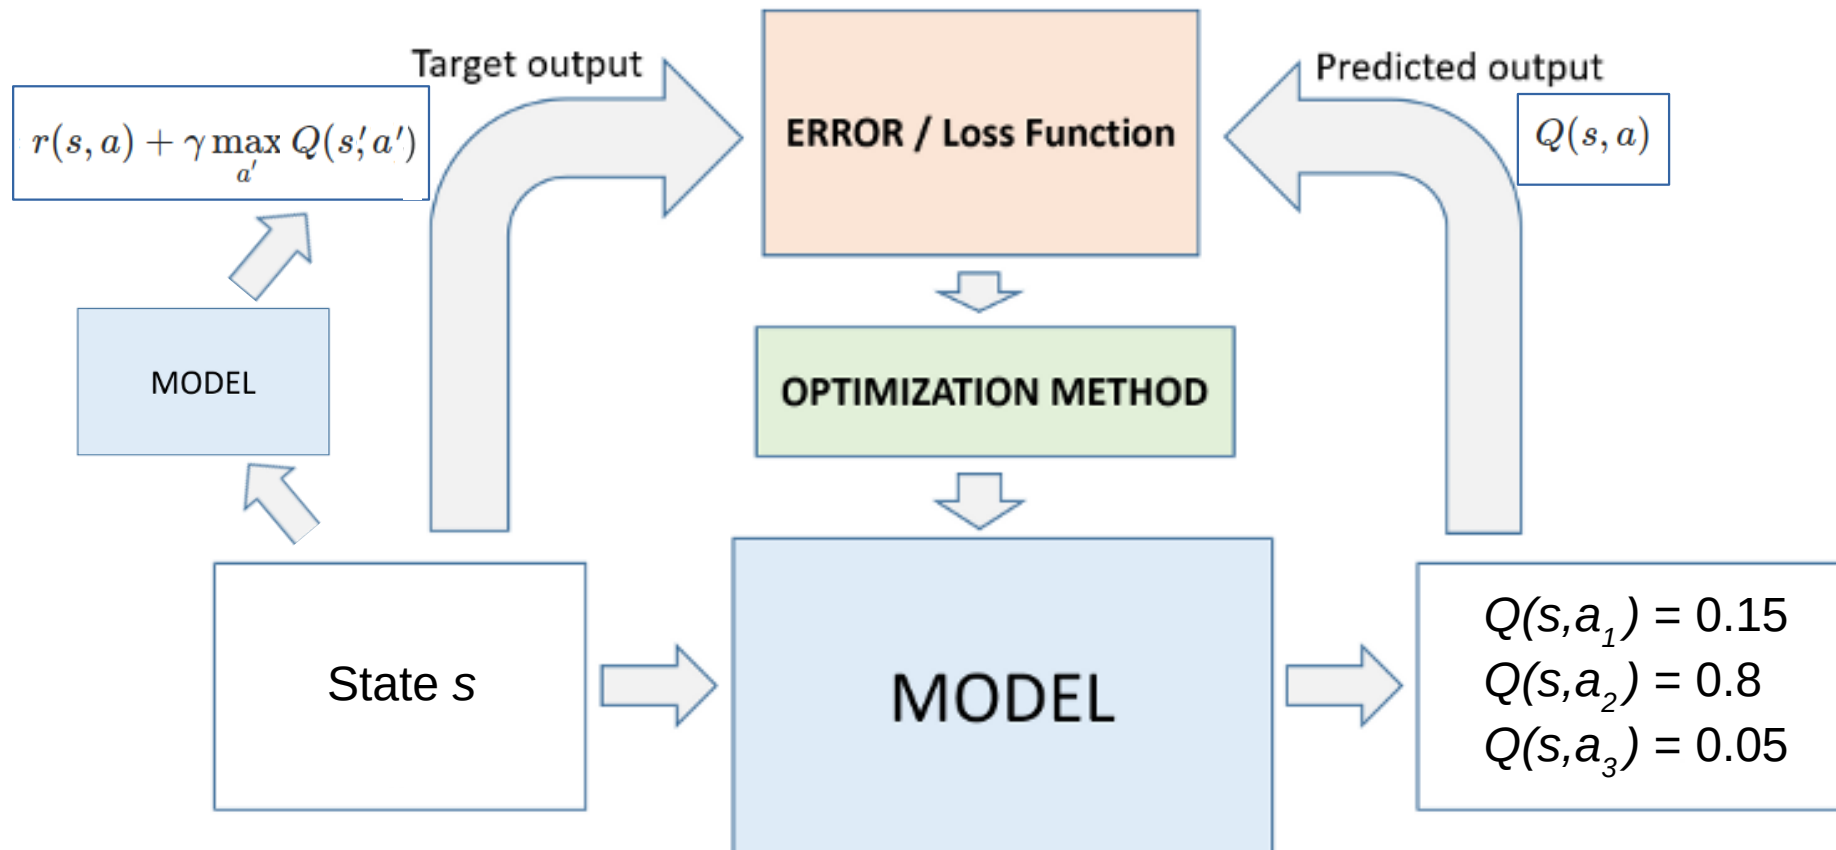
Deep Q-Learning



Deep Q-Learning




Deep Q-Learning




Deep Q-Learning

Traditional supervised learning




Q network (Reinforcement learning)



Deep Q-Learning




Q network (Reinforcement learning)




Deep Q-Learning

Q network (Reinforcement learning)





Prediction 4 →
Prediction 3 →
Prediction 2 →
Prediction 1 →



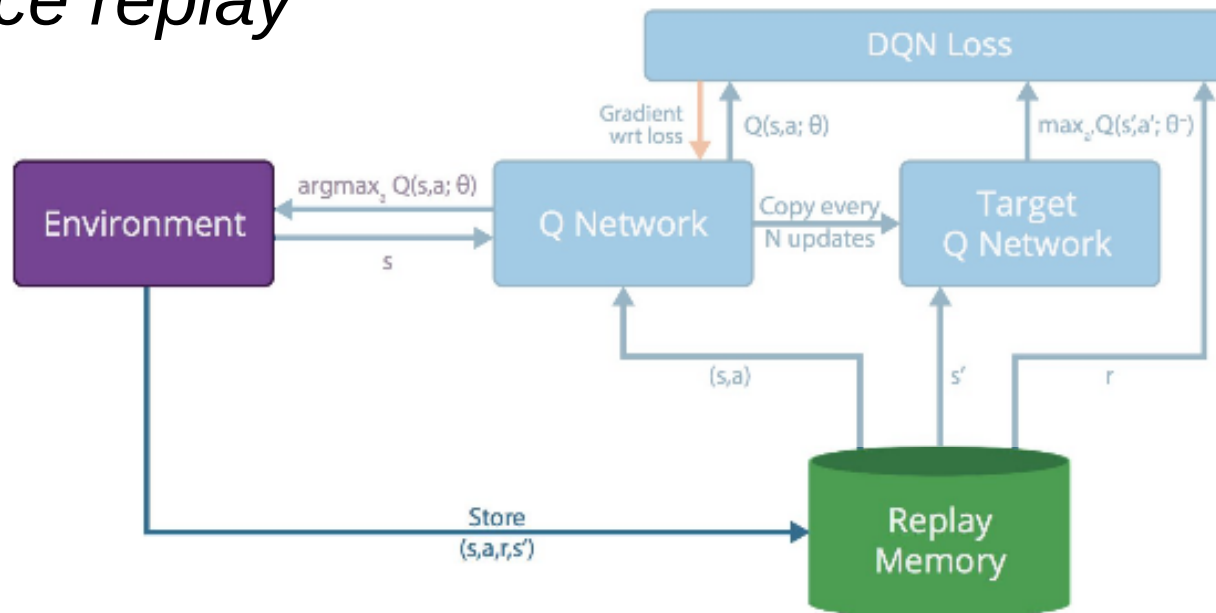
Challenges

- Moving targets
 - *Two networks: behaviour and target*



Challenges

- Moving targets
 - *Two networks: behaviour and target*
- Correlated inputs (due to sequential states)
 - *Experience replay*



Challenges

- Moving targets
 - *Two networks: behaviour and target*
- Correlated inputs (due to sequential states)
 - *Experience replay*
- Markovian
 - *Add recurrency to network*

